

A Quick Introduction to Machine Learning (Scaling)

Lecturer: John Guttag

An Example

1000 patients with 4 features each

Heart rate in beats per minute

Number of past heart attacks

ST segment elevation (binary)


Age



(modified) CC-BY Image Courtesy of Patrick J. Lynch



OGL Image Courtesy of LA(Phot) Stuart Hill/MOD



Binary outcome based on features

Probabilistic, not deterministic

Roughly 31% positive

A Sampling of Examples

P0755:	[48.	1.	0.	62.]:1
P0383:	[103.	1.	1.	99.]:1
P0849:	[42.	1.	1.	92.]:1
P0188:	[71.	2.	0.	58.]:0
P0061:	[87.	1.	0.	79.]:0
P0196:	[52.	0.	0.	85.]:0
P0280:	[78.	0.	0.	81.]:0
P0178:	[50.	1.	0.	59.]:1
P0497:	[80.	0.	0.	58.]:0
P0742:	[78.	2.	0.	72.]:1
P0527:	[78.	1.	0.	60.]:0
P0915:	[60.	2.	0.	57.]:1

Fraction of positives
in total population of
1000 was 0.312

Cluster Using K-means ($k = 3$)

Fraction of positives in population = 0.312

Ran k-means 100 times and chose best clustering

Cluster of size 354 with fraction of positives = 0.338 (1.08x)

Cluster of size 322 with fraction of positives = 0.315 (1.01x)

Cluster of size 324 with fraction of positives = 0.281 (0.90x)

What Happened?

Features have very different means and variance

Heart rate in beats per minute ($\mu = 70$, $\sigma = 15$)

Number of past heart attacks ($\mu = 0.25$, $\sigma = 1.0$)

ST segment elevation (binary)

Age ($\mu = 65$, $\sigma = 15$)

HR and age have higher means & greater dynamic range

Euclidean distance will be biased towards them

Does this seem like a good idea?

“No, No, a Thousand Times No”



Rescale Features

Each variable has same mean and variance

$$x' = \frac{x - \mu_x}{\sigma_x}$$

```
def scaleFeatures(vals):  
    vals = pylab.array(vals)  
    mean = sum(vals)/float(len(vals))  
    sd = stdDev(vals)  
    vals = vals - mean  
    return vals/sd
```

What is the new mean?

What is the new standard deviation?

Testing Scaling

```
def testScaling(n, mean, std):  
    vals = []  
    for i in range(n):  
        vals.append(int(random.gauss(mean, std)))  
    print 'original values', vals  
    sVals = scaleAttrs(vals)  
    print '\n', 'scaled values', sVals  
    print '\n', 'new mean =', sum(sVals)/len(vals)  
    print '\n', 'new sd =', stdDev(sVals)  
  
testScaling(10, 25, 3)
```


Testing Scaling

original values [22, 24, 21, 26, 18, 26, 22, 26,
27, 21]

scaled values [-0.46517657 0.25047969 -0.82300471
0.96613596 -1.89648911 0.96613596
-0.46517657 0.96613596 1.32396409 -0.82300471]

new mean = $-2.22044604925e-16$

new sd = 1.0

The Real Test

Fraction of positives = 0.312

Clustering with unscaled features

Cluster of size 354 with fraction of positives = 0.338 (1.08x)

Cluster of size 322 with fraction of positives = 0.315 (1.01x)

Cluster of size 324 with fraction of positives = 0.281 (0.90x)

Clustering with scaled features

Cluster of size 324 with fraction of positives = 0.055 (0.18x)

Cluster of size 108 with fraction of positives = 0.335 (1.07x)

Cluster of size 568 with fraction of positives = 0.454 (1.45x)